

WESTERN SYDNEY UNIVERSITY



Project Report

Student 1 Chi Hao Lu (19909400)
Student 2 Damian Nguyen (22049939)
Student 3 Qian Heng Yau (22028269)

*Assignment 2 for COMP7023 Predictive Analytics
School of Computer, Data and Mathematical Sciences
Western Sydney University*

Spring 2025

1 Introduction

Global warming has been the hot topic in the World Health Organisation (WHO) for many years. Vehicle emissions are a major source of environmental pollution and climate change, with carbon dioxide (CO₂) and nitrogen oxides (NO_x) being among the most significant pollutants. According to European Environment Agency (2023), CO₂ has impacted the increase of temperature, while NO_x emissions are associated with smog formation, acid rain, and adverse respiratory effects. Hence, the Intergovernmental Panel on Climate Change (IPCC) has suggested to closely monitor the gas emissions from the road vehicles.


This report analyses emissions using the data from the UK Vehicle Certification Agency (VCA), which provides verified measurements of vehicle fuel consumption and emissions under the Worldwide Harmonised Light Vehicle Test Procedure (WLTP). The VCA dataset ensures consistency and accuracy across vehicle testing in accordance with European and international standards (UK Vehicle Certification Agency 2024). Through statistical and predictive modelling, this study explores the key factors influencing emission levels to support sustainable transportation research and evidence-based environmental policymaking.

2 Exploratory Data Analysis

The data set contains of 3702 vehicles from different companies and 44 columns of measurement of vehicle types, fuel consumption and emission.

Category	Columns
Vehicle	Manufacturer, Model, Description, Transmission, Manual or Automatic, Engine Capacity, Fuel Type, Powertrain, Engine Power (PS), Engine Power (Kw)
Fuel Consumption	Diesel VED Supplement, Testing Scheme, WLTP Imperial Low, WLTP Imperial Medium, WLTP Imperial High, WLTP Imperial Extra High, WLTP Imperial Combined, WLTP Imperial Combined (Weighted), WLTP Metric Low, WLTP Metric Medium, WLTP Metric High, WLTP Metric Extra High, WLTP Metric Combined, WLTP Metric Combined (Weighted)
Electric Consumption	Electric energy consumption Miles/kWh, wh/km ,Maximum range (Km), Maximum range (Miles), Equivalent All Electric Range Miles, Equivalent All Electric Range KM, Electric Range City Miles, Electric Range City Km
Standard Emmission	Euro Standard WLTP CO2, WLTP CO2 Weighted, Emissions CO [mg/km], THC Emissions [mg/km], Emissions NOx [mg/km], THC + NOx Emissions [mg/km], Particulates [No.] [mg/km], RDE NOx Urban, RDE NOx Combined, Noise Level dB(A)
Date	Date of change

Out of the 44 columns only 9 columns which are Manufacturer, Model, Description, Transmission, Manual or Automatic, Fuel Type, Powertrain, Euro Standard, and Testing Scheme are the non-numeric data type, the rest of columns are numeric or boolean which can be used modelling. From following chart, we can see the distribution of fuel type of the vehicle in the data set. We can group the vehicles into three main categories: Fuel, Hybrid and Electric.




2.1 Emissions Analysis

As the discussion is about emission, we will further understand the different types of emissions such as Carbon Monoxide (CO), Carbon Dioxide (CO₂), Nitrogen Oxides (NO_x) and Total Hydrocarbons (THC). These emissions are not polluting the environment but impacting our health. According to WHO (2021), when CO binds with haemoglobin more readily than oxygen, it might reduce oxygen transport in blood and posing serious health risks to humans. While CO₂ is non-toxic to humans but it significantly contributes to global warming and climate change as it can trap the heat in the atmosphere (Intergovernmental Panel on Climate Change [IPCC], 2021). When the temperature increases, it might bring disasters to humans such as flood or extreme heat wave.

2.1.1 Carbon Monoxide (CO) and Carbon Dioxide (CO₂)


From the following bar chart, we can clearly see the emissions of CO₂ and CO are co-existent in petrol or hybrid vehicles but not in electric cars. Emission of CO is high especially in diesel vehicles and emission of CO₂ based on WLTP CO₂ is about the same in petrol or diesel vehicles but low in hybrid vehicles. WLTP CO₂ Weighted seems to be not reliable as it is empty for most of the vehicle types.



2.1.2 Nitrogen Oxides (NO_x)


Nitrogen Oxides (NO_x) emissions are a major cause to photochemical smog, acid rain, and respiratory health issues (U.S. Environmental Protection Agency [EPA], 2023). Based on the following chart, we find NO_x emission exists in all the vehicles including electric cars. Electric cars do produce a small average amount of NO_x according to Veza et al (2023). However, it should be understood that the NO_x production from electric vehicles is caused from the source of electricity used to charge the car batteries. NO emissions from electric vehicles originate from the fossil fuel-based electricity used to charge them. The mean for NO_x in Real Driving Emissions (RDE)

for urban and combined are about same but they are higher than Emissions (NO_x) [mg/km] which was the result tested in lab.



2.1.3 Total Hydrocarbons (THC)

Total Hydrocarbons (THC) are volatile organic compounds (VOCs) that react with NO_x in the presence of sunlight to form ground-level ozone and smog (European Environment Agency [EEA], 2022). THC might cause photochemical pollution and impact the urban air quality but it have to react with NO_x, hence, in this report we will not focus in analysing the relationship of THC with other features.



2.2 Target Variables

From the analysis, we have enough data for modelling accross different fuel types but we will choose CO₂ as global warming is most common topic when discussing about the impacts of emission. As mentioned above, WLTP CO₂ Weighted is lacking of data distribution accross

fuel types, hence it will be excluded from target variables. Besides of CO2, we will pick Emissions NOx [mg/km] as our target variables because it has difference mean across fuel type which is easier for our comparison and prediction afterward in this report.

Based on the graphs above, electric cars have no CO2 emissions, and low NOx emissions. Therefore, all of the electric cars will be excluded from the datasets when we build the model. Given the distribution, two datasets will be created based on their fuel types - fuel cars only (data_fuel), and hybrid cars only (data_hybrid). Furthermore, since the target variables (WLTP CO2 and NOx emissions) are not highly included in Electricity /Diesel Fuel type, they will be removed from the data_hybrid.

3 Features Selection

3.1 Model Building

After identified the target variables - Emissions NOx [mg/km] and WLTP CO2, based on the category that we have defined above, we will choose what are the relevant columns to keep in the modelling and remove all irrelevant columns. Reducing the variables/columns can help in improving the performance and cost effectiveness. Variables that are related to the manufacturer or car model will be removed as these variables cannot be used in predictions. As mentioned above, we will create 2 datasets - data_fuel and data_hybrid, which contains fuel vehicles and hybrid vehicles. Hence, the variables that are not relevant to fuel or hybrid will be removed too. For emissions, we will keep only Emissions NOx [mg/km] and WLTP CO2 as our target variables and drop rest of the emissions related columns. Lastly, date and standard category columns will be removed too as this information cannot be used for modeling.


Category	Keep	Remove
Vehicle	Manual or Automatic, Engine Capacity, Fuel Type, Powertrain, Engine Power (PS)	Manufacturer, Model, Description, Transmission, Engine Power (Kw)
Fuel Consumption	WLTP Imperial Low, WLTP Imperial Medium, WLTP Imperial High, WLTP Imperial Extra High, WLTP Imperial Combined, WLTP Imperial Combined (Weighted), WLTP Metric Low, WLTP Metric Medium, WLTP Metric High, WLTP Metric Extra High, WLTP Metric Combined, WLTP Metric Combined (Weighted)	Diesel VED Supplement, Testing Scheme

Category	Keep	Remove
Electric Consumption		Electric energy consumption Miles/kWh, wh/km ,Maximum range (Km), Maximum range (Miles), Equivalent All Electric Range Miles, Equivalent All Electric Range KM, Electric Range City Miles, Electric Range City Km
Standard Emmission	Emissions NOx [mg/km], WLTP CO2	Euro Standard WLTP CO2 Weighted , THC Emissions [mg/km], THC + NOx Emissions [mg/km], Particulates [No.] [mg/km], RDE NOx Urban, RDE NOx Combined, Noise Level dB(A)
Date		Date of change

After removing the irrelevant variables, we have 19 variables remained in the datasets. Next, we will split our data into two groups as mentioned above which are the hybrid vehicles and fuel vehicles. We are getting **1686 hybrid** vehicles and **1250 fuel** vehicles in our datasets after splitting by the fuel type.

3.2 Data Cleaning


Before we start the modelling, cross checking on the data with “not available (NA)” values or empty, this will help to ensure the accuracy of our model. Based on the result, we found NA values in WLTP Imperial Combined (Weighted) and WLTP Metric Combined (Weighted) columns in both fuel and hybrid datasets. After removing the NA rows, we have **1233 fuel** vehicles and **1683 hybrid** vehicles remained in the datasets.




3.3 Correlation Analysis

Correlation analysis started with a heatmap to understand the correlation between each variables. From the heatmap, we can see the correlation with our target variables are low at Engine Power (Ps), WLTP Imperial Combined (Weighted), and WLTP Metric Combined (Weighted).


Correlation Heatmap (lower triangle)



When drilling down to the correlations with CO2, we can see that the WLTP Imperial Combined (Weighted) and WLTP Metric Combined (Weighted) are having the lowest score in both fuel and hybrid vehicles. Whereas fuel is having high correlation with fuel vehicle but not the hybrid vehicle.




In NOx emissions, WLTP Imperial Combined (Weighted) and WLTP Metric Combined (Weighted) are still having lowest correlation in fuel vehicles but these variables are the top three correlated variables in hybrid vehicles. This could be due to the hybrid engine turns on intermittently, it often operates under non-steady, high-load conditions to recharge the battery or accelerate.



3.4 Data Distribution Analysis


Distribution of data is important in identifying the normalisation or standardisation of each variables in our datasets. Below histogram matrix that showing the data distribution for fuel vehicles, we can see that WLTP Imperial Combined (Weighted) and WLTP Metric Combined (Weighted) are not statistically normalised, as their disbutions are tightly clustered and do not follow a standardisation or z-score normalisation. Hence, we can safely remove WLTP Imperial Combined (Weighted) and WLTP Metric Combined (Weighted) columns from *data_fuel* dataset

Histogram Matrix for Fuel Cars



Each variable in hybrid vehicles is statistically normalised, we do not see abnormal data distribution here. Hence, no changes to the variables for hybrid vehicles.

Histogram Matrix for Hyrbid Cars




3.5 Outlier Check

To further improve our model and increase the accuracy, we have perform an analysis of the outliers for both datasets. In the fuel vehicles, we can clearly see the outliers are strongly skewed in most of the WLTP Metric variables and highly skewed in Engine Capacity and Power. This could be due to the high-performance models of vehicles. Since, the WLTP Imperial is having the same purpose with a different unit of measurements, we have decided to drop all the WLTP Metric variables which having higher outliers compare to MLTP Imperial. This table showing the updated model for *data_fuel*.

Category	Keep	Remove
Vehicle	Manual or Automatic, Engine Capacity, Fuel Type, Powertrain, Engine Power (PS)	Manufacturer, Model, Description, Transmission, Engine Power (Kw)

Category	Keep	Remove
Fuel Consumption	WLTP Imperial Low, WLTP Imperial Medium, WLTP Imperial High, WLTP Imperial Extra High, WLTP Imperial Combined	Diesel VED Supplement, Testing Scheme, WLTP Imperial Combined (Weighted), WLTP Metric Combined (Weighted), WLTP Metric Low, WLTP Metric Medium, WLTP Metric High, WLTP Metric Extra High, WLTP Metric Combined
Electric Consumption		Electric energy consumption Miles/kWh, wh/km ,Maximum range (Km), Maximum range (Miles), Equivalent All Electric Range Miles, Equivalent All Electric Range KM, Electric Range City Miles, Electric Range City Km
Standard Emmission	Emissions NOx [mg/km], WLTP CO2	Euro Standard WLTP CO2 Weighted , THC Emissions [mg/km], THC + NOx Emissions [mg/km], Particulates [No.] [mg/km], RDE NOx Urban, RDE NOx Combined, Noise Level dB(A)
Date		Date of change


Outlier Visualisation (Boxplots)



Same for fuel vehicles, it seems to be a high outliers in WLTP Metric for hybrid vehicles. Therefore, these variables will be removed from hybrid vehicles dataset too. Here is the table showing the updated columns remain in *data_hybrid*.

Category	Keep	Remove
Vehicle	Manual or Automatic, Engine Capacity, Fuel Type, Powertrain, Engine Power (PS)	Manufacturer, Model, Description, Transmission, Engine Power (Kw)
Fuel Consumption	WLTP Imperial Low, WLTP Imperial Medium, WLTP Imperial High, WLTP Imperial Extra High, WLTP Imperial Combined	Diesel VED Supplement, Testing Scheme, WLTP Metric Low, WLTP Metric Medium, WLTP Metric High, WLTP Metric Extra High, WLTP Metric Combined, WLTP Metric Combined (Weighted), WLTP Imperial Combined (Weighted)
Electric Consumption		Electric energy consumption Miles/kWh, wh/km ,Maximum range (Km), Maximum range (Miles), Equivalent All Electric Range Miles, Equivalent All Electric Range KM, Electric Range City Miles, Electric Range City Km
Standard		Euro Standard
Emmission	Emissions NOx [mg/km], WLTP CO2	WLTP CO2 Weighted , THC Emissions [mg/km], THC + NOx Emissions [mg/km], Particulates [No.] [mg/km], RDE NOx Urban, RDE NOx Combined, Noise Level dB(A)
Date		Date of change

Outlier Visualisation (Boxplots)



3.5.1 The validation of datasets

```
[1135]: # Data Splitting Function into training set, test set, and validation set
np.random.seed(42) # for reproducibility

def data_split(target: str, data: pd.DataFrame, random_state: int = 42):
    assert target in data.columns, f"Target '{target}' not in dataframe."

    # Split X/y
    X = data.drop(columns=[target])
    y = data[target]

    # Simple numeric imputation (if there are any remaining NaNs)
    # X = X.fillna(X.mean(numeric_only=True)).fillna(0)

    # 80 / 20 split, then 10 / 10 from the 20
    X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.20,
                                                         random_state=random_state)
    X_test, X_val, y_test, y_val = train_test_split(X_temp, y_temp,
                                                    test_size=0.50,
                                                    random_state=random_state)

    # Encode categorical variables as dummy/one-hot columns
    X_train = pd.get_dummies(X_train, drop_first=True, dtype=float)
    X_val = pd.get_dummies(X_val, drop_first=True, dtype=float)
    X_test = pd.get_dummies(X_test, drop_first=True, dtype=float)

    # Re-index columns after encoding
    train_cols = X_train.columns
    X_val = X_val.reindex(columns=train_cols, fill_value=0)
    X_test = X_test.reindex(columns=train_cols, fill_value=0)

    # Scaling the X data
    scaler = StandardScaler()
    X_train = scaler.fit_transform(X_train)
    X_val = scaler.transform(X_val)
    X_test = scaler.transform(X_test)

    return {
        "X_train": X_train,
        "X_val": X_val,
        "X_test": X_test,
        "y_train": y_train,
        "y_val": y_val,
        "y_test": y_test
    }

# reuse this part to access the splitted data.
splits = data_split("WLTP CO2", data_fuel)
```

```

# Access the data
X_train = splits["X_train"]
X_val    = splits["X_val"]
X_test   = splits["X_test"]
y_train  = splits["y_train"]
y_val    = splits["y_val"]
y_test   = splits["y_test"]

print(X_train.shape)
print(X_val.shape)
print(X_test.shape)
print(X_train[0:5,:])

```

```

(986, 13)
(124, 13)
(123, 13)
[[-0.69355105 -0.66468299  0.23878138  0.22261852  0.16192353  0.0135079
   0.15199583 -1.26703709 -0.96022252  0.5139091  -0.19466571  0.12428995
  -0.11558867]
 [ 3.44051473  3.98740879 -2.37420157 -2.1867223  -2.14748521 -2.05174708
  -2.26179734  0.2015991   1.04142527  0.5139091  -0.19466571  0.12428995
  -0.11558867]
 [-0.69355105 -0.50040367  0.59553551  0.43249142  0.31956235  0.087267
   0.29071957 -1.26703709  1.04142527  0.5139091  -0.19466571  0.12428995
  -0.11558867]
 [-0.20296869 -0.23158296 -0.03119472 -0.14675777 -0.18488188 -0.16562136
  -0.08845866 -1.26703709  1.04142527  0.5139091  -0.19466571  0.12428995
  -0.11558867]
 [-0.69453418 -0.81402783  1.51152584  1.47346097  1.34421469  0.86700613
   1.23404104 -0.62849962 -0.96022252  0.5139091  -0.19466571  0.12428995
  -0.11558867]]

```

Fuel car data

```

[1136]: # -- Reusable function for 10-fold CV and evaluation (regression task)--

np.random.seed(42) # for reproducibility

def tenfold_dataValidation(target: str, data: pd.DataFrame, random_state:
    ↪int = 42):
    assert target in data.columns, f"Target '{target}' not in dataframe."

    #reuse this part to access the splitted data.
    splits = data_split(target, data)

    # Access the data
    X_train = splits["X_train"]
    X_val    = splits["X_val"]
    X_test   = splits["X_test"]
    y_train  = splits["y_train"]
    y_val    = splits["y_val"]

```



```

y_test = splits["y_test"]

linear_model = LinearRegression()

# 10-fold CV on TRAIN ONLY (neg MSE -> MSE)
mse_scores = -cross_val_score(linear_model, X_train, y_train, cv=10,
                              scoring="neg_mean_squared_error")

print("10-fold CV MSE scores:", mse_scores.tolist())
print("Mean MSE:", np.mean(mse_scores))

# Auto labeling based on Fuel Type column
if data["Fuel Type"].isin(fuel_types).any():
    data_label = "Fuel Cars"
elif data["Fuel Type"].isin(hybrid_types).any():
    data_label = "Hybrid Cars"
else:
    data_label = "Dataset"

# --- Plot CV MSE ---
plt.figure(figsize=(8, 5))
plt.plot(range(1, 11), mse_scores, marker='o', linestyle='-')
plt.title(f'10-Fold Cross-Validation (MSE) of {target} from_{data_label}')
plt.xlabel('Fold Number'); plt.ylabel('Mean Squared Error (MSE)')
plt.grid(alpha=0.3); plt.xticks(range(1, 11))

# --- Fit on train, check validation ---
linear_model.fit(X_train, y_train)
val_pred = linear_model.predict(X_val)
test_pred = linear_model.predict(X_test)

# MSE on validation and test
print("Validation MSE:", mean_squared_error(y_val, val_pred))
print("Test MSE:", mean_squared_error(y_test, test_pred))
print("\n") #spacing for readability

# RMSE on validation and test
print("Validation RMSE:", np.sqrt(mean_squared_error(y_val, val_pred)))
print("Test RMSE:", np.sqrt(mean_squared_error(y_test, test_pred)))


# return splitted datasets for further use in NN function
# return {
# "X_train": X_train,
# "X_val": X_val,
# "X_test": X_test,
# "y_train": y_train,
# "y_val": y_val,
# "y_test": y_test
# }

```

```
[1137]: # data validation for fuel cars, target variable: WLTP CO2
tenfold_dataValidation("WLTP CO2", data_fuel)
```

10-fold CV MSE scores: [226.8818560077228, 71.07971037711793, 240.
 4021221495773,
 228.03641940795396, 185.59233512324136, 201.05651444852853, 392.8378576500156,
 127.07826439321468, 237.20072334626278, 203.96008070823763]
 Mean MSE: 211.41258836118723
 Validation MSE: 936.3065519621466
 Test MSE: 949.1974967211949

Validation RMSE: 30.599126653585174
 Test RMSE: 30.809048942172733




```
[1138]: # data validation for fuel cars, target variable: Emissions NOx [mg/km]
tenfold_dataValidation("Emissions NOx [mg/km]", data_fuel)
```

10-fold CV MSE scores: [196.263888695339, 166.97656776826773,
 171.84763097495428, 236.30575475101722, 164.63380342084642, 171.
 42723437305426,
 193.74103894193985, 585.9913886695763, 154.96234536980253, 197.3585544035697]
 Mean MSE: 223.95082073683676
 Validation MSE: 1400.8457844798568
 Test MSE: 1293.3702178748638

Validation RMSE: 37.42787443176351

Test RMSE: 35.96345670086322




Hybrid car data

```
[1139]: # data validation for hybrid cars, target variable: WLTP CO2  
tenfold_dataValidation("WLTP CO2", data_hybrid)
```

10-fold CV MSE scores: [1239.2532655144914, 1586.0682053838518,
1366.9974995948303, 1170.7755565874538, 1552.5254169525424, 1825.
9652911355797,
1507.1878461500476, 1271.4742172473284, 970.3356191917294, 1391.3434076879907]
Mean MSE: 1388.1926325445847
Validation MSE: 1036.097747870503
Test MSE: 1539.9931032653658

Validation RMSE: 32.18847228233274

Test RMSE: 39.242745868062876



```
[1140]: # data validation for hybrid cars, target variable: Emissions NOx [mg/km]
tenfold_dataValidation("Emissions NOx [mg/km]", data_hybrid)
```

10-fold CV MSE scores: [42.27947119632712, 47.956806391015, 66.14828240354515, 57.94611387863636, 33.30093578990461, 42.873359949927355, 42.90273724034652, 36.751902710119815, 56.835284641779026, 39.163033607400436]


Mean MSE: 46.615792780900144

Validation MSE: 71.62983677331998

Test MSE: 24.3941498342779

Validation RMSE: 8.463441189806897

Test RMSE: 4.939043412876415



3.6 Lasso

```
[1141]: def feature_select_lasso(
    X_train, X_test, y_train, y_test, *,
    X_val=None, y_val=None,                    # <-- NEW (optional)
    validation_mode="merge",                  # "merge" or "separate"
    refit_after_cv=True,                      # only used when
    validation_mode="separate"
    features=None,
    dataset_name="Dataset",
    cv=10,
    random_state=42,
    max_iter=5000,
    show_plots=True
):
    """
    Lasso feature selection with 10-fold CV + learning curve.
    Supports 3-way split via `validation_mode`:
        - "merge": CV+fit on train+val, test on test (recommended)
        - "separate": CV on train only; report val metrics; optionally refit on
    train+val
    """

    # ----- choose columns -----
    use_cols = X_train.columns.tolist() if features is None else list(dict.
    fromkeys(features))
```

```

for c in use_cols:
    if c not in X_train.columns: raise ValueError(f"Missing column in_
↪X_train: {c}")
Xtr, Xte = X_train[use_cols].copy(), X_test[use_cols].copy()
Xva = None if X_val is None else X_val[use_cols].copy()

# ----- preprocessing -----
num_cols = Xtr.select_dtypes(include=["float64", "int64"]).columns.
↪tolist()
cat_cols = [c for c in use_cols if c not in num_cols]
pre = ColumnTransformer([
    ("num", Pipeline([("imp", SimpleImputer(strategy="mean")),
                      ("sc", StandardScaler())]), num_cols),
    ("cat", "drop", cat_cols)
])

# ----- choose fitting data -----
if validation_mode == "merge" and X_val is not None and y_val is not_
↪None:
    Xfit = pd.concat([Xtr, Xva], axis=0)
    yfit = pd.concat([y_train, y_val], axis=0)
    fit_label = "train+val"
else:
    Xfit, yfit = Xtr, y_train
    fit_label = "train"

# ----- CV model on fitting data -----
cv_model = LassoCV(cv=cv, random_state=random_state, max_iter=max_iter)
pipe_cv = Pipeline([("pre", pre), ("model", cv_model)])
pipe_cv.fit(Xfit, yfit)
alpha_star = pipe_cv.named_steps["model"].alpha_

# ----- optional refit (separate mode) -----
final_pipe = pipe_cv
if validation_mode == "separate" and refit_after_cv and X_val is not_
↪None and y_val is not None:
    # lock alpha and refit on train+val for the fairest test estimate
    final_model = Lasso(alpha=alpha_star, random_state=random_state,
↪max_iter=max_iter)
    final_pipe = Pipeline([("pre", pre), ("model", final_model)])
    final_pipe.fit(pd.concat([Xtr, Xva]), pd.concat([y_train, y_val]))
    fit_label = "train+val (refit with fixed )"

# ----- predictions -----
yhat_tr = final_pipe.predict(Xtr)
yhat_te = final_pipe.predict(Xte)
yhat_va = None if Xva is None else final_pipe.predict(Xva)

# ----- metrics (RSS + MSE) -----
rss_train = float(np.sum((y_train - yhat_tr)**2))

```

```

mse_train = float(mean_squared_error(y_train, yhat_tr))
rss_test  = float(np.sum((y_test - yhat_te)**2))
mse_test  = float(mean_squared_error(y_test, yhat_te))

rss_val = mse_val = None
if Xva is not None and y_val is not None:
    rss_val = float(np.sum((y_val - yhat_va)**2))
    mse_val = float(mean_squared_error(y_val, yhat_va))

print(f"\n===== {dataset_name.upper()}")
print(f"Features ({len(use_cols)}): {use_cols}")
print(f"Validation mode: {validation_mode} | Fitted on: {fit_label}")
print(f"Chosen alpha ( ) from CV: {alpha_star:.6f}")
if mse_val is not None:
    print(f"VAL RSS: {rss_val:.2f} | VAL MSE: {mse_val:.2f}")
print(f"TRAIN RSS: {rss_train:.2f} | TRAIN MSE: {mse_train:.2f}")
print(f"TEST RSS: {rss_test:.2f} | TEST MSE: {mse_test:.2f}")

# ----- overfit check -----
ratio = rss_test / rss_train if rss_train>0 else np.inf
gap    = (mse_test - mse_train)/mse_train*100 if mse_train>0 else np.inf
status = "OK"
if ratio>2 or gap>50: status="Possible overfit"
elif ratio<0.6 and gap<-30: status="Well-regularised / easy test split"
print(f"Overfit check → RSS ratio {ratio:.2f}, MSE gap {gap:.1f}% → {status}")

# ----- learning curve (on the data used for CV) -----
train_sizes, tr_scores, va_scores = learning_curve(
    pipe_cv, Xfit, yfit, cv=cv, scoring="neg_mean_squared_error",
    train_sizes=np.linspace(0.1, 1.0, 8), random_state=random_state
)
lc_train = -tr_scores.mean(axis=1)
lc_val    = -va_scores.mean(axis=1)

# ----- plots -----
if show_plots:
    fig, axs = plt.subplots(2, 2, figsize=(12,8))
    fig.suptitle(f"{dataset_name} - LassoCV Diagnostics", fontsize=14,
fontweight="bold")

    # CV curve from the CV model
    axs[0,0].semilogx(pipe_cv.named_steps["model"].alphas_,
        pipe_cv.named_steps["model"].mse_path_.
mean(axis=1), marker='.')
    axs[0,0].axvline(alpha_star, ls='--', color='r')
    axs[0,0].set_title(f"CV MSE vs Alpha ({cv}-fold)")
    axs[0,0].set_xlabel("Alpha ( )"); axs[0,0].set_ylabel("Mean CV MSE")

```

```

    # Predicted vs Actual (test)
    axs[0,1].scatter(y_test, yhat_te)
    lo, hi = float(np.min(y_test)), float(np.max(y_test))
    axs[0,1].plot([lo, hi], [lo, hi], 'r--')
    axs[0,1].set_title("Predicted vs Actual (Test)")
    axs[0,1].set_xlabel("Actual WLTP CO2"); axs[0,1].
↪set_ylabel("Predicted WLTP CO2")

    # Learning curve (on Xfit/yfit)
    axs[1,0].plot(train_sizes, lc_train, 'o-', label="Train MSE")
    axs[1,0].plot(train_sizes, lc_val, 's-', label="Validation MSE")
    axs[1,0].set_title(f"Learning Curve on {fit_label}")
    axs[1,0].set_xlabel("Training set size"); axs[1,0].set_ylabel("MSE");
↪ axs[1,0].legend()

    # Coefficients
    coef = final_pipe.named_steps["model"].coef_
    names = final_pipe.named_steps["pre"].get_feature_names_out()
    axs[1,1].barh(names, coef); axs[1,1].set_title("Lasso Coefficients")
    axs[1,1].set_xlabel("Coefficient")
    plt.tight_layout(); plt.show()

    # ----- return -----
    return {
        "dataset": dataset_name,
        "features": use_cols,
        "alpha": float(alpha_star),
        "mode": validation_mode,
        "fitted_on": fit_label,
        "rss_train": rss_train, "mse_train": mse_train,
        "rss_val": rss_val, "mse_val": mse_val,
        "rss_test": rss_test, "mse_test": mse_test,
        "overfit_flag": status,
        "learning_curve": {"train_sizes": train_sizes, "train_mse": ↪
↪lc_train, "val_mse": lc_val}
    }, final_pipe

```

```

[1142]: # Example 1 - Fuel: manually pick columns
(X_train, X_val, X_test, y_train, y_val, y_test) = [
    data_split("WLTP CO2", data_fuel)[k] for k in
    ["X_train", "X_val", "X_test", "y_train", "y_val", "y_test"]
]
fuel_feats = ['Manual or Automatic', 'Engine Capacity', 'Engine Power (PS)', ↪
↪'WLTP Metric Low', 'WLTP Metric Medium', 'WLTP Metric High', 'WLTP Metric ↪
↪Extra High', 'WLTP Metric Combined', 'Emissions NOx [mg/km]']
#res_fuel_manual, pipe_fuel_manual = feature_select_lasso(
#    X_train, X_test, y_train, y_test,
#    X_val=X_val, y_val=y_val,
#    validation_mode="separate",
#    features=fuel_feats, cv=10, show_plots=True

```



```
#)
```

```
[1143]: # Example 1 - Fuel: manually pick columns
(X_train, X_val, X_test, y_train, y_val, y_test) = [
    data_split("WLTP CO2", data_fuel)[k] for k in
    ["X_train", "X_val", "X_test", "y_train", "y_val", "y_test"]
]
fuel_feats = ['Engine Capacity', 'Engine Power (PS)', 'WLTP Imperial_
↳ Combined']
#res_fuel_manual, pipe_fuel_manual = feature_select_lasso(
#    X_train, X_test, y_train, y_test,
#    X_val=X_val, y_val=y_val,
#    validation_mode="separate",
#    features=fuel_feats, cv=10, show_plots=True
#)
```

```
[1144]: # Example 2 - Hybrid: use all columns (features=None)
(X_train, X_val, X_test, y_train, y_val, y_test) = [
    data_split("WLTP CO2", data_hybrid)[k] for k in
    ["X_train", "X_val", "X_test", "y_train", "y_val", "y_test"]
]

#res_hybrid_all, pipe_hybrid_all = feature_select_lasso(
#    X_train, X_test, y_train, y_test,
#    X_val=X_val, y_val=y_val,
#    validation_mode="merge",
#    features=None, cv=10, show_plots=True
#)
```

```
[1145]: # Example 3 - Lock alpha to a known value (e.g., from a previous CV run)
(X_train, X_val, X_test, y_train, y_val, y_test) = [
    data_split("WLTP CO2", data_hybrid)[k] for k in
    ["X_train", "X_val", "X_test", "y_train", "y_val", "y_test"]
]

#res_fuel_locked, pipe_fuel_locked = feature_select_lasso(
#    X_train, X_test, y_train, y_test,
#    X_val=X_val, y_val=y_val,
#    validation_mode="merge",
#    features=fuel_feats, cv=10, show_plots=True
#)
```

3.7 Neural Network - Regression Tasks

Variables Selection The following variables should be used to develop neural networks model.

```
[1146]: #print current variables from the fuel cars dataset
print(data_fuel.columns)
print(data_fuel.shape)
```

```
Index(['Manual or Automatic', 'Engine Capacity', 'Fuel Type', 'Powertrain',
      'Engine Power (PS)', 'WLTP Imperial Low', 'WLTP Imperial Medium',
      'WLTP Imperial High', 'WLTP Imperial Extra High',
      'WLTP Imperial Combined', 'WLTP CO2', 'Emissions NOx [mg/km]'],
      dtype='object')
(1233, 12)
```

3.7.1 Neural Network Regression (2 Layers)

```
[1147]: # Reusable function for Neural Network regression task - simple 1 hidden
        ↪ layer NN
def neural_network_regression(target: str, data: pd.DataFrame, random_state:
        ↪ int = 42):
    assert target in data.columns, f"Target '{target}' not in dataframe."

    #reuse this part to access the splitted data.
    splits = data_split(target, data)

    # Access the data
    X_train = splits["X_train"]
    X_val    = splits["X_val"]
    X_test   = splits["X_test"]
    y_train  = splits["y_train"]
    y_val    = splits["y_val"]
    y_test   = splits["y_test"]

    # Flatten target arrays
    y_train = y_train.values.reshape(-1, 1)
    y_val    = y_val.values.reshape(-1, 1)
    y_test   = y_test.values.reshape(-1, 1)

    # Define NN dimensions
    N = len(y_train) # number of samples
    D_in = X_train.shape[1] # number of features - 12
    D_out = 1 # number of outputs
    H = 20 # number of hidden neurons, based on number of features

    # Initialize weights
    W0 = normal(0, 0.1, size=(D_in, H))
    W1 = normal(0, 0.1, size=(H, D_out))

    # Training loop
    eta = 0.0005

    loss_history = []
    for i in range(2001):
        # forward pass

        ## Bias Terms
        b0 = np.zeros((1, H))
```

```

b1 = np.zeros((1, 1))
# z1 = X.dot(W0) + b0
# a2 = a1.dot(W1) + b1

z1 = X_train.dot(W0) + b0
z1 = np.clip(z1, -500, 500)
a1 = 1 / (1 + np.exp(-z1))
a2 = a1.dot(W1) + b1
Loss = np.mean((a2 - y_train) ** 2)
loss_history.append(Loss)

# # print loss every 100 iterations
# if i % 100 == 0:
#     print(f"Iteration {i}: Loss {Loss:.4f}")

# backward pass
da2 = 2.0 * (a2 - y_train)
dW1 = a1.T.dot(da2) / N
da1 = da2.dot(W1.T)
dz1 = da1 * a1 * (1 - a1)
dW0 = X_train.T.dot(dz1) / N

# update
W0 -= eta * dW0 # weights update for features
W1 -= eta * dW1

# Evaluate
z1 = X_train.dot(W0)
z1 = np.clip(z1, -500, 500)
a1 = 1 / (1 + np.exp(-z1))
y_pred = a1.dot(W1)

# MSE and RMSE on training data
print('Training RMSE:', np.sqrt(mean_squared_error(y_train, y_pred)))
print('Coefficient of determination ( $R^2$ ) on training data:', □)
↪round(r2_score(y_train, y_pred),4))

# Evaluate on test set
z1 = X_test.dot(W0)
z1 = np.clip(z1, -500, 500)
a1 = 1 / (1 + np.exp(-z1))
y_test_pred = a1.dot(W1)

print('Test RMSE:', np.sqrt(mean_squared_error(y_test, y_test_pred)))
print('Coefficient of determination ( $R^2$ ) on test data:', □)
↪round(r2_score(y_test, y_test_pred),4))

# Evaluate on validation set
z1 = X_val.dot(W0)
z1 = np.clip(z1, -500, 500)

```

```

a1 = 1 / (1 + np.exp(-z1))
y_val_pred = a1.dot(W1)
print('Validation RMSE:', np.sqrt(mean_squared_error(y_val, y_val_pred)))
print('Coefficient of determination (R²) on validation data:',
round(r2_score(y_val, y_val_pred),4))

# Auto labeling based on Fuel Type column
if data["Fuel Type"].isin(fuel_types).any():
    data_label = "Fuel Cars"
elif data["Fuel Type"].isin(hybrid_types).any():
    data_label = "Hybrid Cars"
else:
    data_label = "Dataset"

#Visualise the MSE
plt.figure(figsize=(8,5))
plt.plot(loss_history, color='teal')
plt.title(f"Neural Network Training Loss of {target} from {data_label}\n (MSE over iterations)",
        fontsize = 15, pad = 10)
plt.xlabel("Iteration")
plt.ylabel("Mean Squared Error (MSE)")
plt.grid(alpha=0.3)
plt.show()

# # Return the weights and features of the model
# return{
#     "features": features,
#     "W0": W0
# }


# Neural Network regression for fuel cars, target variable: WLTP CO2
neural_network_regression("WLTP CO2", data_fuel)

```

```

Training RMSE: 11.83214538566079
Coefficient of determination (R²) on training data: 0.9541
Test RMSE: 43.92853084963061
Coefficient of determination (R²) on test data: 0.1397
Validation RMSE: 57.19261359806857
Coefficient of determination (R²) on validation data: 0.1014

```



```
[1148]: # Neural Network regression for fuel cars, target variable: Emissions NOx ↴  
        ↵ [mg/km]  
        neural_network_regression("Emissions NOx [mg/km]", data_fuel)
```

Training RMSE: 14.821462889091457

Coefficient of determination (R^2) on training data: 0.1043


Test RMSE: 11.84380327307547

Coefficient of determination (R^2) on test data: 0.0628

Validation RMSE: 12.432183269388801

Coefficient of determination (R^2) on validation data: 0.0942

Neural Network Training Loss of Emissions NOx [mg/km] from Fuel Cars
(MSE over iterations)



```
[1149]: # Neural Network regression for hybrid cars, target variable: WLTP CO2  
neural_network_regression("WLTP CO2", data_hybrid)
```

Training RMSE: 38.45094043133957


Coefficient of determination (R^2) on training data: 0.4958

Test RMSE: 38.71394959730473

Coefficient of determination (R^2) on test data: 0.4605

Validation RMSE: 30.797058493922243

Coefficient of determination (R^2) on validation data: 0.5314



```
[1150]: # Neural Network regression for hybrid cars, target variable: Emissions NOx [mg/km]
neural_network_regression("Emissions NOx [mg/km]", data_hybrid)
```

Training RMSE: 6.951501409137495

Coefficient of determination (R^2) on training data: 0.358


Test RMSE: 4.555392989431061

Coefficient of determination (R^2) on test data: 0.5398

Validation RMSE: 8.564928331625294

Coefficient of determination (R^2) on validation data: 0.2628

Neural Network Training Loss of Emissions NOx [mg/km] from Hybrid Cars
(MSE over iterations)



```
[1151]: # #Examine weights from the NNR
# import io, contextlib

# buf = io.StringIO()
# with contextlib.redirect_stdout(buf): # hide print outputs
#     res = neural_network_regression("WLTP CO2", data_fuel);

# # from your Regressor() output
# W0 = res["W0"] # shape: [n_features, H]
# features = res["features"]
# imp = np.mean(np.abs(W0), axis=1)

# print("W0 shape:", W0.shape)
# print("Number of features:", len(features))


# # Ensure they match in length
# min_len = min(len(features), len(imp))
# features = features[:min_len]
# imp = imp[:min_len]

# imp_df = pd.DataFrame({"Feature": features, "Importance": imp})
# imp_df = imp_df.sort_values(by="Importance", ascending=True)

[1152]: # importance = average absolute incoming weight to hidden layer
plt.figure(figsize=(8, 10))
#plt.barh(imp_df["Feature"], imp_df["Importance"])
plt.title("Input → Hidden Weight Magnitudes (avg |W0|)")
```



```
plt.xlabel("Average |weight|")  
plt.tight_layout()  
plt.show()
```



3.7.2 Neural Network Regression - 3 Layers

```
[1153]: # Reusable function for Neural Network regression task - 2 hidden layer NN
def neural_network_regression_2hidden(target: str, data: pd.DataFrame,
    random_state: int = 42):
    assert target in data.columns, f"Target '{target}' not in dataframe."

    splits = data_split(target, data)

    X_train = splits["X_train"]
    X_val = splits["X_val"]
    X_test = splits["X_test"]
    y_train = splits["y_train"]
    y_val = splits["y_val"]
    y_test = splits["y_test"]

    # Flatten target arrays
    y_train = y_train.values.reshape(-1, 1)
    y_val = y_val.values.reshape(-1, 1)
    y_test = y_test.values.reshape(-1, 1)

    # scaler = StandardScaler()
    # X_train_scaled = scaler.fit_transform(X_train)
    # X_test_scaled = scaler.transform(X_test)
    # X_val_scaled = scaler.transform(X_val)

    # Define NN dimensions
    N = len(y_train) # number of samples
    D_in = X_train.shape[1] # number of features
    D_out = 1 # number of outputs
    H1 = 30 # number of hidden neurons in hidden layer 1
    H2 = 20 # number of hidden neurons in hidden layer 2

    # -- Initialize weights
    # Layer 1: input -> hidden1
    W0 = normal(0, 0.1, size=(D_in, H1))
    b0 = np.zeros((1, H1))

    # Layer 2: hidden1 -> hidden2
    W1 = normal(0, 0.1, size=(H1, H2))
    b1 = np.zeros((1, H2))

    # Output: hidden2 -> output
    W2 = normal(0, 0.1, size=(H2, D_out))
    b2 = np.zeros((1, D_out))

    # Training loop
    eta = 0.0003

    loss_history = []
    for i in range(3001):
```

```

# ---- forward pass ----
## Layer 1
z1 = X_train.dot(W0) + b0
z1 = np.clip(z1, -500, 500)
a1 = 1.0 / (1.0 + np.exp(-z1)) # sigmoid

# Layer 2
z2 = a1.dot(W1) + b1
z2 = np.clip(z2, -500, 500)
a2 = 1.0 / (1.0 + np.exp(-z2)) # sigmoid

# Output (linear)
y_hat = a2.dot(W2) + b2

# Loss (MSE)
Loss = np.mean((y_hat - y_train) ** 2)
loss_history.append(float(Loss))

# ---- backward pass ----
# dLoss/dy_hat = 2*(y_hat - y) / N
dy = (2.0 / N) * (y_hat - y_train)

# Gradients for output layer (a2 -> y_hat)
dW2 = a2.T.dot(dy)
db2 = np.sum(dy, axis=0, keepdims=True)

# Backprop into hidden2
da2 = dy.dot(W2.T)
dz2 = da2 * a2 * (1.0 - a2)
dW1 = a1.T.dot(dz2)
db1 = np.sum(dz2, axis=0, keepdims=True)

# Backprop into hidden1
da1 = dz2.dot(W1.T)
dz1 = da1 * a1 * (1.0 - a1)
dW0 = X_train.T.dot(dz1)
db0 = np.sum(dz1, axis=0, keepdims=True)

# ---- update parameters ----
W2 -= eta * dW2
b2 -= eta * db2
W1 -= eta * dW1
b1 -= eta * db1
W0 -= eta * dW0
b0 -= eta * db0

# # Evaluate
# z1 = X_train_scaled.dot(W0)
# z1 = np.clip(z1, -500, 500)
# a1 = 1 / (1 + np.exp(-z1))
# y_pred = a1.dot(W2)

```

```

# Evaluate on training set
z1_ = X_train.dot(W0) + b0
z1_ = np.clip(z1_, -500, 500)
a1_ = 1.0 / (1.0 + np.exp(-z1_))

z2_ = a1_.dot(W1) + b1
z2_ = np.clip(z2_, -500, 500)
a2_ = 1.0 / (1.0 + np.exp(-z2_))

y_pred = a2_.dot(W2) + b2

# MSE and RMSE on training data
print('Training RMSE:', np.sqrt(mean_squared_error(y_train, y_pred)))
print('Coefficient of determination ( $R^2$ ) on training data:',
round(r2_score(y_train, y_pred),4))

# Evaluate on test set
z1_ = X_test.dot(W0) + b0
z1_ = np.clip(z1_, -500, 500)
a1_ = 1.0 / (1.0 + np.exp(-z1_))
z2_ = a1_.dot(W1) + b1
z2_ = np.clip(z2_, -500, 500)
a2_ = 1.0 / (1.0 + np.exp(-z2_))
y_test_pred = a2_.dot(W2) + b2

print('Test RMSE:', np.sqrt(mean_squared_error(y_test, y_test_pred)))
print('Coefficient of determination ( $R^2$ ) on test data:',
round(r2_score(y_test, y_test_pred),4))

# Evaluate on validation set
z1_ = X_val.dot(W0) + b0
z1_ = np.clip(z1_, -500, 500)
a1_ = 1.0 / (1.0 + np.exp(-z1_))
z2_ = a1_.dot(W1) + b1
z2_ = np.clip(z2_, -500, 500)
a2_ = 1.0 / (1.0 + np.exp(-z2_))
y_val_pred = a2_.dot(W2) + b2

print('Validation RMSE:', np.sqrt(mean_squared_error(y_val, y_val_pred)))
print('Coefficient of determination ( $R^2$ ) on validation data:',
round(r2_score(y_val, y_val_pred),4))

# Auto labeling based on Fuel Type column
if data["Fuel Type"].isin(fuel_types).any():
    data_label = "Fuel Cars"
elif data["Fuel Type"].isin(hybrid_types).any():
    data_label = "Hybrid Cars"
else:

```

```

data_label = "Dataset"

#Visualise the MSE
plt.figure(figsize=(8,5))
plt.plot(loss_history, color='teal')
plt.title(f"Neural Network (3 Layer) Training Loss of \n {target} from_{data_label}",
↪fontstyle=15, pad=10)
plt.xlabel("Iteration")
plt.ylabel("Mean Squared Error (MSE)")
plt.grid(alpha=0.3)
plt.show()

# # Return the weights and features of the model
# return{
#     "features": features,
#     "WO": WO
# }

# Neural Network regression for fuel cars, target variable: WLTP CO2
neural_network_regression_2hidden("WLTP CO2", data_fuel)

```

Training RMSE: 15.622564437897724


Coefficient of determination (R^2) on training data: 0.92

Test RMSE: 15.850467231513218

Coefficient of determination (R^2) on test data: 0.888

Validation RMSE: 18.690134370355924

Coefficient of determination (R^2) on validation data: 0.904



```
[1154]: # NN Regression (2 hidden layers), Emissions NOx from Fuel cars
neural_network_regression_2hidden("Emissions NOx [mg/km]", data_fuel)
```

Training RMSE: 15.393607895371947


Coefficient of determination (R^2) on training data: 0.0338

Test RMSE: 12.28769373326525

Coefficient of determination (R^2) on test data: -0.0088

Validation RMSE: 12.82621176054791

Coefficient of determination (R^2) on validation data: 0.0358



```
[1155]: # NN Regression (2 hidden layers), CO2 emissions from Hybrid cars
neural_network_regression_2hidden("WLTP CO2", data_hybrid)
```

Training RMSE: 35.42839109030671


Coefficient of determination (R^2) on training data: 0.5719

Test RMSE: 32.62246511055373

Coefficient of determination (R^2) on test data: 0.6169

Validation RMSE: 32.63765510527334

Coefficient of determination (R^2) on validation data: 0.4737



```
[1156]: # NN Regression (2 hidden layers), Emissions NOx from Hybrid cars
neural_network_regression_2hidden("Emissions NOx [mg/km]", data_hybrid)
```

Training RMSE: 8.005424250010037


Coefficient of determination (R^2) on training data: 0.1486

Test RMSE: 5.914581063728446

Coefficient of determination (R^2) on test data: 0.2243

Validation RMSE: 9.415363072073598

Coefficient of determination (R^2) on validation data: 0.1092



3.7.3 Neural Network Regression - 4 Layers

[1157]: *# Reusable function for Neural Network regression task - 3 hidden layer NN*

```
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def neural_network_regression_3hidden(target: str, data: pd.DataFrame,
    random_state: int = 42):
    assert target in data.columns, f"Target '{target}' not in dataframe."

    # Perform data splitting
    splits = data_split(target, data)

    X_train = splits["X_train"]
    X_val = splits["X_val"]
    X_test = splits["X_test"]
    y_train = splits["y_train"]
    y_val = splits["y_val"]
    y_test = splits["y_test"]

    # Flatten target arrays
    y_train = y_train.values.reshape(-1, 1)
    y_val = y_val.values.reshape(-1, 1)
    y_test = y_test.values.reshape(-1, 1)
```



```

# Define NN dimensions
N = len(y_train) # number of samples
D_in = X_train.shape[1] # number of features
D_out = 1 # number of outputs

H1 = 40 # number of hidden neurons in hidden layer 1
H2 = 30 # number of hidden neurons in hidden layer 2
H3 = 20 # number of hidden neurons in hidden layer 3

# -- Initialize weights
# Layer 1: input -> hidden1
W0 = normal(0, 0.1, size=(D_in, H1))
b0 = np.zeros((1, H1))

# Layer 2: hidden1 -> hidden2
W1 = normal(0, 0.1, size=(H1, H2))
b1 = np.zeros((1, H2))

# Layer 3: hidden2 -> hidden3
W2 = normal(0, 0.1, size=(H2, H3))
b2 = np.zeros((1, H3))

# Output: hidden3 -> output
W3 = normal(0, 0.1, size=(H3, D_out))
b3 = np.zeros((1, D_out))

# Training loop
eta = 0.0005

loss_history = []
for i in range(2001):
    # ---- forward pass ----
    ## Layer 1
    z1 = X_train.dot(W0) + b0
    z1 = np.clip(z1, -500, 500)
    a1 = 1 / (1 + np.exp(-z1)) # sigmoid

    # Layer 2
    z2 = a1.dot(W1) + b1
    z2 = np.clip(z2, -500, 500)
    a2 = 1 / (1 + np.exp(-z2)) # sigmoid

    # Layer 3
    z3 = a2.dot(W2) + b2
    z3 = np.clip(z3, -500, 500)
    a3 = 1 / (1 + np.exp(-z3)) # sigmoid

    # Output (linear)
    y_hat = a3.dot(W3) + b3

```

```

# Loss (MSE)
Loss = np.mean((y_hat - y_train) ** 2)
loss_history.append(float(Loss))

# ---- backward pass ----
dy = (2.0 / N) * (y_hat - y_train)

# Gradients for output layer (a3 -> y_hat)
dW3 = a3.T.dot(dy)
db3 = np.sum(dy, axis=0, keepdims=True)

# Backprop into hidden3
da3 = dy.dot(W3.T)
dz3 = da3 * a3 * (1 - a3)
dW2 = a2.T.dot(dz3)
db2 = np.sum(dz3, axis=0, keepdims=True)

# Backprop into hidden2
da2 = dz3.dot(W2.T)
dz2 = da2 * a2 * (1 - a2)
dW1 = a1.T.dot(dz2)
db1 = np.sum(dz2, axis=0, keepdims=True)

# Backprop into hidden1
da1 = dz2.dot(W1.T)
dz1 = da1 * a1 * (1 - a1)
dW0 = X_train.T.dot(dz1)
db0 = np.sum(dz1, axis=0, keepdims=True)

# ---- update parameters ----
W3 -= eta * dW3
b3 -= eta * db3

W2 -= eta * dW2
b2 -= eta * db2

W1 -= eta * dW1
b1 -= eta * db1

W0 -= eta * dW0
b0 -= eta * db0

# Evaluate on training set
def predict(X):
    z1 = X.dot(W0) + b0
    z1_ = np.clip(z1, -500, 500)
    a1 = 1.0 / (1.0 + np.exp(-z1_))

    z2 = a1.dot(W1) + b1

```

```

    z2_ = np.clip(z2, -500, 500)
    a2 = np.clip(z2_, -500, 500)

    z3 = a2.dot(W2) + b2
    z3_ = np.clip(z3, -500, 500)
    a3 = 1.0 / (1.0 + np.exp(-z3_))

    return a3.dot(W3) + b3

# z1 = X_train_scaled.dot(W0) + b0
# z1_ = np.clip(z1, -500, 500)
# a1 = 1.0 / (1.0 + np.exp(-z1_))

# z2 = a1.dot(W1) + b1
# z2_ = np.clip(z2, -500, 500)
# a2 = np.clip(z2_, -500, 500)

# z3 = a2.dot(W2) + b2
# z3_ = np.clip(z3, -500, 500)
# a3 = 1.0 / (1.0 + np.exp(-z3_))
# y_train_pred = a3.dot(W3) + b3

# MSE and RMSE on training data
y_train_pred = predict(X_train)
print('Training RMSE:', np.sqrt(mean_squared_error(y_train,
↪y_train_pred)))
print('Coefficient of determination (R²) on training data:',
↪round(r2_score(y_train, y_train_pred),4))

# Evaluate on test set
y_test_pred = predict(X_test)
print('Test RMSE:', np.sqrt(mean_squared_error(y_test, y_test_pred)))
print('Coefficient of determination (R²) on test data:',
↪round(r2_score(y_test, y_test_pred),4))

# Evaluate on validation set
y_val_pred = predict(X_val)
print('Validation RMSE:', np.sqrt(mean_squared_error(y_val, y_val_pred)))
print('Coefficient of determination (R²) on validation data:',
↪round(r2_score(y_val, y_val_pred),4))

# Auto labeling based on Fuel Type column
if data["Fuel Type"].isin(fuel_types).any():
    data_label = "Fuel Cars"
elif data["Fuel Type"].isin(hybrid_types).any():
    data_label = "Hybrid Cars"
else:
    data_label = "Dataset"

```

```


#Visualise the MSE
plt.figure(figsize=(8,5))
plt.plot(loss_history, color='teal')
plt.title(f"Neural Network (2 Layer) Training Loss of {target} from_{data_label} \n (MSE over iterations)",
          fontsize = 15, pad = 10)
plt.xlabel("Iteration")
plt.ylabel("Mean Squared Error (MSE)")
plt.grid(alpha=0.3)
plt.show()

# # Return the weights and features of the model
# return{
#     "features": features,
#     "WO": WO
# }

# Neural Network regression for fuel cars, target variable: WLTP CO2
neural_network_regression_3hidden("WLTP CO2", data_fuel)

```


Training RMSE: 138.1645823541116
 Coefficient of determination (R^2) on training data: -5.2572
 Test RMSE: 138.67875772052057
 Coefficient of determination (R^2) on test data: -7.5738
 Validation RMSE: 146.49132268610106
 Coefficient of determination (R^2) on validation data: -4.8955



```
[1158]: neural_network_regression_3hidden("Emissions NOx [mg/km]", data_fuel)
```

Training RMSE: 15.601972032883996
Coefficient of determination (R^2) on training data: 0.0075
Test RMSE: 12.140231683389397
Coefficient of determination (R^2) on test data: 0.0153
Validation RMSE: 12.977859002448692
Coefficient of determination (R^2) on validation data: 0.0129


Neural Network (2 Layer) Training Loss of Emissions NOx [mg/km] from Fuel Cars
(MSE over iterations)



```
[1159]: neural_network_regression_3hidden("WLTP CO2", data_hybrid)
```

Training RMSE: 35.56435477602406
Coefficient of determination (R^2) on training data: 0.5687
Test RMSE: 32.2759250197809
Coefficient of determination (R^2) on test data: 0.625
Validation RMSE: 32.67101399065757
Coefficient of determination (R^2) on validation data: 0.4726

Neural Network (2 Layer) Training Loss of WLTP CO2 from Hybrid Cars
(MSE over iterations)



```
[1160]: neural_network_regression_3hidden("Emissions NOx [mg/km]", data_hybrid)
```

Training RMSE: 8.549650805680743

Coefficient of determination (R^2) on training data: 0.0289


Test RMSE: 6.486216803484162

Coefficient of determination (R^2) on test data: 0.0671

Validation RMSE: 9.968267522196147

Coefficient of determination (R^2) on validation data: 0.0015

Neural Network (2 Layer) Training Loss of Emissions NOx [mg/km] from Hybrid Cars
(MSE over iterations)



4 References

- European Environment Agency, 2023, ‘Greenhouse Gas Emissions from Transport in Europe.’, *EEA Report No. 23/2023*, Publications Office of the European Union.
- Intergovernmental Panel on Climate Change (IPCC), 2022, ‘Climate Change 2022: Mitigation of Climate Change’, *Contribution of Working Group III to the Sixth Assessment Report of the IPCC*, Cambridge University Press.
- UK Vehicle Certification Agency, 2024, *Fuel Consumption and Emission Data Tables: Car Fuel and CO Information. Department for Transport*, viewed 29 September 2025, <https://www.vehicle-certification-agency.gov.uk/>
- World Health Organization, 2021, *WHO Global Air Quality Guidelines: Particulate Matter, Ozone, Nitrogen Dioxide, Sulfur Dioxide and Carbon Monoxide*, Geneva: WHO.
- U.S. Environmental Protection Agency (EPA), 2023, *Nitrogen Oxides (NO_x) Pollution*, EPA.
- European Environment Agency (EEA), 2022, *Air Pollution Sources — Volatile Organic Compounds (VOCs)*, EEA Report
- Veza, Ibharm et al, 2023, ‘Electric Vehicle (EV) and Driving Towards Sustainability: Comparison between EV, HEV, PHEV, and ICE Vehicles to Achieve Net Zero Emissions by 2050 from EV’, *Alexandria Engineering Journal*, vol. 82, pp. 459–467, DOI.10.1016/j.aej.2023.10.020